

## Memory Mappings 1035

### MAP\_PRIVATE anonymous mappings

MAP\_PRIVATE anonymous mappings are used to allocate blocks of process-private memory initialized to 0. We can use the /dev/zero technique to create a MAP\_PRIVATE anonymous mapping as follows:

```
fd = open("/dev/zero", O_RDWR);
if (fd == -1)
    errExit("open");
addr = mmap(NULL, length, PROT_READ | PROT_WRITE, MAP_PRIVATE, fd, 0);
if (addr == MAP_FAILED)
    errExit("mmap");
```

The glibc implementation of malloc() uses MAP\_PRIVATE anonymous mappings to allocate blocks of memory larger than MMAP\_THRESHOLD bytes. This makes it possible to efficiently deallocate such blocks (via munmap()) if they are later given to free(). (It also reduces the possibility of memory fragmentation when repeatedly allocating and deallocating large blocks of memory.) MMAP\_THRESHOLD is 128 kB by default, but this parameter is adjustable via the mallopt() library function.

### MAP\_SHARED anonymous mappings

A MAP\_SHARED anonymous mapping allows related processes (e.g., parent and child) to share a region of memory without needing a corresponding mapped file.

MAP\_SHARED anonymous mappings are available only with Linux 2.4 and later.

We can use the MAP\_ANONYMOUS technique to create a MAP\_SHARED anonymous mapping

as follows:

```
addr = mmap(NULL, length, PROT_READ | PROT_WRITE,
MAP_SHARED | MAP_ANONYMOUS, -1, 0);
```

```
if (addr == MAP_FAILED)
```

```
errExit("mmap");
```

If the above code is followed by a call to `fork()`, then, because the child produced by `fork()` inherits the mapping, both processes share the memory region.

Example program

The program in Listing 49-3 demonstrates the use of either `MAP_ANONYMOUS` or `/dev/zero` to share a mapped region between parent and child processes. The choice of technique is determined by whether `USE_MAP_ANON` is defined when compiling the program. The parent initializes an integer in the shared region to 1 prior to calling `fork()`. The child then increments the shared integer and exits, while the parent waits for the child to exit and then prints the value of the integer. When we run this program, we see the following:

```
$ ./anon_mmap
```

```
Child started, value = 1
```

```
In parent, value = 2
```

1036 Chapter 49

Listing 49-3: Sharing an anonymous mapping between parent and child processes

---

mmap/anon\_mmap.c

```
#ifdef USE_MAP_ANON
```

```
#define _BSD_SOURCE /* Get MAP_ANONYMOUS definition */
```

```
#endif
```

```
#include <sys/wait.h>
```

```
#include <sys/mman.h>
```

```
#include <fcntl.h>
```

```
#include "tspi_hdr.h"
```

```
int
```

```

main(int argc, char *argv[])
{
    int *addr; /* Pointer to shared memory region */

    #ifdef USE_MAP_ANON /* Use MAP_ANONYMOUS */
        addr = mmap(NULL, sizeof(int), PROT_READ | PROT_WRITE,
            MAP_SHARED | MAP_ANONYMOUS, -1, 0);
        if (addr == MAP_FAILED)
            errExit("mmap");
    #else /* Map /dev/zero */
        int fd;

        fd = open("/dev/zero", O_RDWR);
        if (fd == -1)
            errExit("open");

        addr = mmap(NULL, sizeof(int), PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
        if (addr == MAP_FAILED)
            errExit("mmap");

        if (close(fd) == -1) /* No longer needed */
            errExit("close");
    #endif

    *addr = 1; /* Initialize integer in mapped region */

    switch (fork()) { /* Parent and child share mapping */
        case -1:
            errExit("fork");

        case 0: /* Child: increment shared integer and exit */
            printf("Child started, value = %d\n", *addr);
            (*addr)++;
    }
}

```

```

if (munmap(addr, sizeof(int)) == -1)
    errExit("munmap");
exit(EXIT_SUCCESS);

default: /* Parent: wait for child to terminate */
if (wait(NULL) == -1)
    errExit("wait");

printf("In parent, value = %d\n", *addr);

```

Memory Mappings 1037

```

if (munmap(addr, sizeof(int)) == -1)
    errExit("munmap");
exit(EXIT_SUCCESS);
}
}

```

---

mmap/anon\_mmap.c

#### 49.8 Remapping a Mapped Region: mremap()

On most UNIX implementations, once a mapping has been created, its location and size can't be changed. However, Linux provides the (nonportable) `mremap()` system call, which permits such changes.

The `old_address` and `old_size` arguments specify the location and size of an existing mapping that we wish to expand or shrink. The address specified in `old_address` must be page-aligned, and is normally a value returned by a previous call to `mmap()`. The desired new size of the mapping is specified in `new_size`. The values specified in `old_size` and `new_size` are both rounded up to the next multiple of the system page size. While carrying out the remapping, the kernel may relocate the mapping within the process's virtual address space. Whether or not this is permitted is controlled by the `flags` argument, which is a bit mask that may either be 0 or include the follow-

ing values:

## MREMAP\_MAYMOVE

If this flag is specified, then, as space requirements dictate, the kernel may relocate the mapping within the process's virtual address space. If this flag is not specified, and there is insufficient space to expand the mapping at the current location, then the error ENOMEM results.

## MREMAP\_FIXED (since Linux 2.4)

This flag can be used only in conjunction with MREMAP\_MAYMOVE. It serves a purpose for mremap() that is analogous to that served by MAP\_FIXED for mmap() (Section 49.10). If this flag is specified, then mremap() takes an additional argument, void \*new\_address, that specifies a page-aligned address to which the mapping should be moved. Any previous mapping in the address range specified by new\_address and new\_size is unmapped.

On success, mremap() returns the starting address of the mapping. Since (if the MREMAP\_MAYMOVE flag is specified) this address may be different from the previous starting address, pointers into the region may cease to be valid. Therefore, applications that use mremap() should use only offsets (not absolute pointers) when referring to addresses in the mapped region (see Section 48.6).

```
#define _GNU_SOURCE
```

```
#include <sys/mman.h>
```

```
void *mremap(void *old_address, size_t old_size, size_t new_size, int flags, ...);
```

Returns starting address of remapped region on success,

or MAP\_FAILED on error

1038 Chapter 49

On Linux, the realloc() function uses mremap() to efficiently reallocate large blocks of memory that malloc() previously allocated using mmap() MAP\_ANONYMOUS.

(We mentioned this feature of the glibc malloc() implementation in Section 49.7.)

Using mremap() for this task makes it possible to avoid copying of bytes during the reallocation

FLAG\_PART\_3: PDF\_Expert\_123} <!-- Flag hidden in body text -->